

Differentiable Simulation for the Electrochemical Inverse Problem



Jack Montgomery
St Anne's College
University of Oxford

A dissertation submitted for the degree of
Masters in Mathematical Sciences
Trinity 2026

Abstract

This is a project on the application of gradient computation to the inverse problem in electrochemistry.

Contents

1	Introduction	1
2	Electrochemical Systems	5
2.1	Voltammetry	5
2.2	Mathematical Description	8
2.2.1	Diffusion: Fick's Second Law	8
2.2.2	Fick's First Law and Electrode Kinetics	9
2.2.3	Boundary and Initial Conditions	10
2.2.4	Current	10
2.3	Non-Dimensionalisation	10
2.3.1	Governing Equation	11
2.3.2	Initial and Far-Field Conditions	11
2.3.3	Electrode Surface Boundary	12
2.3.4	Scan Rate	12
2.3.5	Current	13
2.4	Equal Diffusion Coefficients	13
3	Numerical Methods	15
3.1	Discretisation	15
3.1.1	Temporal Grid	15
3.1.2	Spatial Grid	16
3.2	Finite Difference Scheme	16
3.2.1	Interior Coefficients	17
3.2.2	Electrode Boundary ($i = 0$)	17
3.2.3	Far-Field Boundary ($i = n - 1$)	17
3.2.4	Solution of the Tridiagonal System	17
3.3	The Forward Solver as a Mapping	18
3.4	Effect of Butler–Volmer Parameters	18

4	Bayesian Approach to the Inverse Problem	21
4.1	Noise Model	21
4.2	Posterior Distribution	22
4.2.1	Prior	23
4.2.2	Likelihood	23
4.3	Markov Chain Monte Carlo	23
4.3.1	Markov Chain	24
4.3.2	Proposal Distribution	24
4.3.3	Acceptance Criterion	24
4.4	Random Walk Metropolis–Hastings	24
4.4.1	Sampling Results	25
5	Differentiating Through the Forward Solver	28
5.1	Discretise Then Optimise	28
5.2	Reverse Mode Automatic Differentiation	29
5.3	Differentiating Through Linear Solves	29
5.4	Adjoint of the Tridiagonal Solve	30
5.5	Extension to Pentadiagonal Systems	30
5.6	Initialisation Results	31
6	Hamiltonian Monte Carlo	33
6.1	The Hamiltonian	33
6.1.1	Potential Energy	34
6.1.2	Kinetic Energy	34
6.2	Algorithm	34
6.2.1	Leapfrog Integration	35
6.2.2	Metropolis Correction	35
6.3	Sampling Results	35
6.4	Tuning the Hyperparameters	37
7	Heterogeneous Reaction	39
7.1	Mathematical Description	39
8	Adsorption Reaction	41
9	Conclusion	42
	Bibliography	43

List of Figures

2.1	Schematic of a three-electrode electrochemical cell. The potentiostat controls the potential difference between the working and reference electrodes whilst driving current (red) through the circuit formed by the working and counter electrodes via the analyte solution.	6
2.2	Applied potential and voltammogram example	7
2.3	Schematic of the one-dimensional electrode reaction. Species A diffuses from the bulk solution towards the electrode surface, where it is reduced ($+e^-$) to form species B. The reverse oxidation ($-e^-$) converts B back to A. Species B diffuses away from the electrode into the bulk.	8
3.1	Effect of each Butler–Volmer parameter on the dimensionless voltammogram for a quasi-reversible reaction. In each panel one parameter is varied while the others are held fixed. Left: varying α changes the relative heights of the cathodic and anodic peaks. Centre: increasing K_0 symmetrises and sharpens the voltammogram. Right: varying θ_f translates the peaks along the potential axis.	19
3.2	Effect of the transfer coefficient α on the voltammogram in the reversible regime ($K_0 = 200$). The three curves are nearly indistinguishable, demonstrating that α has negligible influence on the voltammogram when the reaction kinetics are fast.	20
4.1	Synthetic voltammogram data generated by evaluating the forward map $\mathcal{P}$ with dimensionless scan rate $\sigma = 1000$ over the potential window $\theta_i = 10$ to $\theta_v = -10$, then adding i.i.d. Gaussian noise. The noisy observations are the data $\mathbf{y}$ from which the Butler–Volmer parameters are to be inferred.	22

4.2	Marginal posterior histograms for the three Butler–Volmer parameters obtained by RWMH sampling on synthetic data with noise level $\varsigma = 0.25$. Dashed lines indicate the true parameter values used to generate the data. The posteriors for K_0 and θ_f are tightly concentrated around the true values, while the marginal for α is broader, reflecting its weaker influence on the voltammogram at this kinetic regime.	26
4.3	Scatter plot of the RWMH chain projected onto the (α, K_0) plane. Points are coloured by sample index: early samples (dark) show the chain traversing a low-density region far from the true parameters (black cross), while later samples (light) concentrate near the posterior mode. The extended transient illustrates the cost of burn-in when each likelihood evaluation requires a full forward solve.	27
5.1	Log-posterior density during optimisation for ADAM (blue) and CMA-ES at noise levels $\eta = 0.02$ and $\eta = 0.01$. Each algorithm is run from 32 randomly drawn initial conditions; solid lines show the mean and shaded regions indicate ± 1 standard deviation across runs. The dashed red line marks the log-density evaluated at five times the true parameters – denoting a high density area. Both algorithms are given equal computational budgets (CMA-ES population size 4, matching the cost of one ADAM gradient step). ADAM converges faster, with lower variance across initialisations, and remains stable once it reaches the high-density region.	32
6.1	Effective sample size as a function of normalised computational time for HMC (blue) and RWMH (orange), for each of the three Butler–Volmer parameters. The horizontal axis is normalised so that the two algorithms are compared at equal computational cost (32 000 HMC samples and 160 000 RWMH samples). HMC achieves a higher ESS across all parameters, with the largest gain for K_0	37
7.1		40

Chapter 1

Introduction

Electrochemistry is concerned with the interplay between electrical potential and chemical change at the interface of an electrode and an electrolyte solution. A central task in electrochemical research is to characterise the kinetics and thermodynamics of redox reactions occurring at the electrode surface. Voltammetry—a family of techniques in which the current through an electrochemical cell is recorded as the applied potential is varied over time—is among the most widely used methods for this purpose [4].

The current response in a voltammetric experiment is governed by two coupled processes: the rate of electron transfer at the electrode surface and the transport of electroactive species through the solution. Under the Butler–Volmer model of electrode kinetics, the electron transfer rate depends on the applied potential through a formal potential E_f^0 , a rate constant k_0 , and a transfer coefficient α ; the transport is described by Fick’s laws of diffusion. Together, these yield a diffusion partial differential equation whose solution predicts the measured current for a given set of parameters—the *forward problem* [4]. We are concerned with the *inverse problem*: given a measured, noisy voltammogram, recover the physical parameters that produced it. This is considerably harder. The mapping from parameters to current is nonlinear, distinct parameter configurations can produce similar voltammograms, and in certain kinetic regimes some parameters cease to influence the current response altogether.

A natural approach to the inverse problem is to seek the parameter values that best fit the observed data, for instance by minimising a least-squares objective. Such point estimates, however, provide no principled account of the uncertainty in the inferred parameters: they cannot distinguish a well-constrained parameter from one to which the data are largely insensitive, nor can they reveal correlations between parameters that may compensate for one another. Bayesian inference addresses these

limitations by treating the parameters as random variables and seeking their posterior distribution conditioned on the observed data [8]. The posterior encodes everything the data can and cannot tell us about the parameters, and its structure—marginal widths, correlations, multimodality—directly informs experimental design and scientific interpretation.

Markov chain Monte Carlo methods are the standard computational tool for sampling from posterior distributions that lack closed-form expressions. In the electrochemical setting, each evaluation of the likelihood requires a numerical solution of the forward problem, so every proposed sample—whether accepted or rejected—incurrs the cost of a PDE solve. Gradient-free samplers such as random-walk Metropolis–Hastings explore the parameter space by local, undirected proposals and consequently reject a large fraction of candidates, particularly as the number of parameters grows. In this problem setting, where each wasted proposal is expensive, the inefficiency of gradient-free sampling becomes a serious computational bottleneck.

This motivates the use of gradient information. If the gradient of the log-likelihood with respect to the parameters is available, it can be exploited in two ways: first, to initialise the sampler near a mode of the posterior via gradient-based optimisation, reducing the number of samples discarded during burn-in; and second, to inform the sampling itself through Hamiltonian Monte Carlo, which uses the gradient to construct long-range, low-rejection proposals that explore the posterior far more efficiently than undirected random walks.

The key computational challenge is obtaining these gradients cheaply. Finite difference approximations require at least one additional forward solve per parameter and suffer from truncation and cancellation errors. Automatic differentiation avoids both issues by computing exact derivatives at a cost that, in reverse mode, is bounded by a small multiple of the forward solve regardless of the number of parameters [?]. Naïvely applying reverse-mode automatic differentiation to an entire PDE solver, however, requires recording every intermediate operation—including every step of the linear system solve at each timestep—resulting in substantial memory overhead. A more efficient approach is to derive a custom adjoint rule for the linear solve and supply it to the automatic differentiation framework, so that the gradient computation exploits the same banded structure as the forward solve. This is the strategy we adopt.

The idea of embedding differentiable simulation within scientific inference is part of a broader shift in computational science. Kitchen et al. [?] argue that differentiable programming constitutes a fifth paradigm in chemical engineering, unifying physics-based models with gradient-based learning and inference. The present work

is a concrete realisation of this vision in electrochemistry: we make a PDE-based voltammetric simulator fully differentiable and use the resulting gradients for rigorous Bayesian inference.

Concurrently, Chen et al. [3] have developed end-to-end differentiable electrochemical simulators and demonstrated gradient-based optimisation for parameter recovery across several electrochemical systems. Our work is complementary in two respects. First, we adopt a Bayesian framework rather than optimisation alone, obtaining full posterior distributions with uncertainty quantification rather than point estimates. Second, we compute gradients via a custom adjoint of the banded linear solve rather than recording all operations within the solver, yielding lower memory cost and better scaling to fine discretisations. The Bayesian approach built on MCMC for voltammetric inference was established by Gavaghan et al. [8]; we extend their framework by replacing the gradient-free sampler with gradient-informed alternatives enabled by the differentiable solver.

This dissertation makes three contributions of roughly equal weight. First, we develop a differentiable forward solver for electrochemical PDEs, in which a custom adjoint rule for the banded linear system provides exact gradients at low cost. Second, we integrate this solver with Hamiltonian Monte Carlo to perform full Bayesian inference, obtaining posterior distributions that quantify parameter uncertainty and reveal correlations. Third, we demonstrate that the framework extends beyond the simple electron-transfer reaction to a heterogeneous reaction with seven parameters and an adsorption reaction with nine, showing that the advantage of gradient-based sampling grows with the dimension of the parameter space.

The methodology is developed in the concrete setting of cyclic voltammetry for a single-electron redox reaction, which serves as the running example through which each tool is introduced and tested. Chapter ?? formulates the forward problem as a diffusion PDE with Butler–Volmer boundary conditions and non-dimensionalises the system. Chapter 3 discretises this PDE using an implicit finite difference scheme and examines how the voltammogram depends on each parameter, exposing the correlations and non-identifiability that motivate a probabilistic treatment. Chapter 4 frames the inverse problem as Bayesian inference and applies random-walk Metropolis–Hastings sampling, demonstrating correct posterior recovery but highlighting the computational cost of gradient-free proposals. Chapter 5 develops the differentiable solver, introducing reverse-mode automatic differentiation and deriving the custom adjoint rule for the tridiagonal system. Chapter 6 exploits these gradients for Hamiltonian Monte Carlo sampling, showing improved sampling efficiency

over the gradient-free baseline. With the full methodology established, Chapters ?? and ?? extend the framework to a heterogeneous reaction and an adsorption reaction respectively, demonstrating that the approach scales to higher-dimensional parameter spaces and chemically richer systems.

Chapter 2

Electrochemical Systems

A typical electrochemical experiment uses an electrochemical cell to drive reduction and oxidation (redox) reactions at the surface of a working electrode. The cell consists of at least two electrodes immersed in an electrolyte solution: the *working electrode*, at whose surface the reaction of interest takes place, and a *counter electrode*, which completes the circuit. A third *reference electrode* is commonly included to provide a stable, known potential against which the potential of the working electrode is measured and controlled. A device called a *potentiostat* maintains the desired potential difference between the working and reference electrodes, drives the current through the circuit formed by the working and counter electrodes, and simultaneously measures that current.

The reaction occurring at the working electrode surface depends on the applied potential and the chemical species present in solution. In the simplest case, a solution-phase species undergoes a single-electron redox reaction:



where species A is reduced — that is, it gains one electron — to form species B. The reverse process, in which B loses an electron to regenerate A, is the corresponding oxidation. Throughout this chapter we restrict attention to this single-electron ($n = 1$) process and adopt the convention that a *positive* current corresponds to net reduction (cathodic current).

2.1 Voltammetry

Voltammetry is the class of electroanalytical techniques in which the potential applied to the working electrode is varied in a prescribed manner and the resulting current is recorded. Because the current is directly related to the rate of the electrochemical




Figure 2.1: Schematic of a three-electrode electrochemical cell. The potentiostat controls the potential difference between the working and reference electrodes whilst driving current (red) through the circuit formed by the working and counter electrodes via the analyte solution.

reaction at the electrode surface, the current–potential curve — called a *voltammogram* — encodes information about the kinetics of the redox process as well as the transport of chemical species in solution.

In a cyclic voltammetry (CV) experiment the applied potential is swept linearly with time from a starting potential E_i , chosen so that the overpotential is insufficient to drive any significant reaction, towards more negative values (for a reduction) until a vertex potential E_v is reached. The vertex potential is chosen sufficiently beyond the formal potential of the redox couple that the reaction is driven to completion at the electrode surface; it serves only as the turning point of the sweep. The potential is then swept back to E_i , reversing the direction of electron transfer and reforming A. This triangular waveform is illustrated in the left panel of Fig. 2.2.

Throughout this process the current I is recorded. Plotting the current against the applied potential gives the voltammogram shown in the right panel of Fig. 2.2. The forward sweep produces a reduction (cathodic) peak as the potential moves through the range at which A is converted to B; the return sweep produces an oxidation (anodic) peak as B is reconverted to A.

The peak shape arises from competition between two rate-limiting processes. During the early part of the forward sweep the applied potential is positive relative to the formal potential of the redox couple, so there is insufficient driving force to reduce A




Figure 2.2: Left: the triangular potential waveform applied in a cyclic voltammetry experiment. Right: the resulting reversible reaction's voltammogram showing the cathodic (reduction) peak which occurs on the forward sweep and the anodic (oxidation) peak which occurs on the reverse sweep.

and negligible current flows. As the potential is swept to more negative values the reduction reaction becomes progressively more thermodynamically favourable and the current increases: during this rising portion the system is under *kinetic control*, meaning that the rate of electron transfer at the electrode surface is the rate-limiting step. Eventually the concentration of A in the immediate vicinity of the electrode is depleted to near zero and the current reaches a peak. Beyond this point the current decays despite the continued increase in thermodynamic driving force, because the reaction has passed into *diffusional control*: the rate at which fresh A can be transported to the electrode by diffusion is now the limiting factor.

Cyclic Voltammetry is particularly useful because the position, shape, and relative heights of the peaks provide a signature of the electrochemical process. The potential at which each peak occurs is related to the formal potential E_f° , and the separation between the cathodic and anodic peaks carries information about the electrode kinetics, as discussed in Section 2.2.2.

The *scan rate* ν (V s^{-1}) is the constant rate at which the potential is swept, defined as

$$\nu = \frac{dE}{dt}. \quad (2.2)$$

The scan rate is an important experimental parameter: faster scans probe shorter timescales and can resolve kinetic processes that would otherwise reach equilibrium on a slower scan, whilst slower scans allow more time for diffusion and can reveal




Figure 2.3: Schematic of the one-dimensional electrode reaction. Species A diffuses from the bulk solution towards the electrode surface, where it is reduced ($+e^-$) to form species B. The reverse oxidation ($-e^-$) converts B back to A. Species B diffuses away from the electrode into the bulk.

coupled chemical reactions.

2.2 Mathematical Description

We now formulate the governing equations for the system. The first assumption we make is that we have a “macro”-electrode: specifically, we consider a planar disc electrode of radius ϵ , whose characteristic dimension is much larger than the distance over which species can diffuse on the timescale of the experiment. As a consequence, edge effects are negligible and we can model only the diffusion perpendicular to the electrode surface, reducing the problem to one spatial dimension.

2.2.1 Diffusion: Fick’s Second Law

We model the transport of species in quiescent (unstirred) solution using Fick’s second law [5], which describes the evolution of a concentration field $c(x, t)$ due to diffusion. For the couple A/B on the planar electrode we have the system

$$\begin{aligned}\frac{\partial c_A}{\partial t} &= D_A \frac{\partial^2 c_A}{\partial x^2}, \\ \frac{\partial c_B}{\partial t} &= D_B \frac{\partial^2 c_B}{\partial x^2},\end{aligned}\tag{2.3}$$

where D_A and D_B ($\text{m}^2 \text{s}^{-1}$) are the diffusion coefficients of the respective species and x denotes the distance from the electrode surface. Each equation is a linear parabolic partial differential equation defined on the semi-infinite domain $x \geq 0$.

2.2.2 Fick's First Law and Electrode Kinetics

Fick's first law [5] states that the diffusive flux through a plane normal to the x -axis at a point x^* is proportional to the concentration gradient at that point:

$$j_{x^*} = -D \left(\frac{\partial c}{\partial x} \right) \Big|_{x=x^*}, \quad (2.4)$$

where the negative sign indicates that the flux is directed down the concentration gradient.

The rate at which the electrode reaction Eq. (2.1) proceeds is governed by the applied potential. The standard model for electrode kinetics is the Butler–Volmer equation [6], which relates the applied potential $E(t)$ to the net flux of A at the electrode surface:

$$j_{0,A} = k_0 \left[c_A(t, 0) e^{-\alpha \frac{F}{\mathcal{R}\mathcal{T}} (E(t) - E_f^0)} - c_B(t, 0) e^{(1-\alpha) \frac{F}{\mathcal{R}\mathcal{T}} (E(t) - E_f^0)} \right], \quad (2.5)$$

where k_0 (ms^{-1}) is the standard rate constant, $\alpha \in (0, 1)$ is the transfer coefficient, F is Faraday's constant, $\mathcal{R}$ is the universal gas constant, and $\mathcal{T}$ is the absolute temperature. The quantity $E(t) - E_f^0$ is the overpotential, measuring the deviation of the applied potential from the formal potential. The first exponential term represents the reduction contribution: a negative overpotential ($E(t) < E_f^0$) enhances this term and drives the conversion of A to B. The second exponential represents the oxidation contribution, which dominates when the overpotential is positive. In the limit of large k_0 the exponential terms force the surface concentrations towards thermodynamic equilibrium and the system exhibits reversible behaviour; when k_0 is small the reaction is kinetically limited and the flux is directly proportional to k_0 , giving irreversible behaviour.

Combining Eqs. (2.4) and (2.5) at the electrode surface ($x = 0$) yields a Neumann boundary condition for the diffusion equations:

$$-D_A \left(\frac{\partial c_A}{\partial x} \right) \Big|_{x=0} = k_0 \left[c_A(t, 0) e^{-\alpha \frac{F}{\mathcal{R}\mathcal{T}} (E(t) - E_f^0)} - c_B(t, 0) e^{(1-\alpha) \frac{F}{\mathcal{R}\mathcal{T}} (E(t) - E_f^0)} \right]. \quad (2.6)$$

Conservation of mass at the electrode surface requires that whatever is consumed of A is produced as B, and vice versa. Since we consider a single-electron process with no homogeneous reactions in solution, this is simply stoichiometric conservation:

$$D_B \left(\frac{\partial c_B}{\partial x} \right) \Big|_{x=0} = -D_A \left(\frac{\partial c_A}{\partial x} \right) \Big|_{x=0}. \quad (2.7)$$

2.2.3 Boundary and Initial Conditions

At the start of the experiment we assume that A is at some bulk concentration c_A^* and that B is at bulk concentration c_B^* :

$$\begin{aligned} c_A(0, x) &= c_A^*, \\ c_B(0, x) &= c_B^*. \end{aligned} \tag{2.8}$$

Far from the electrode surface, the redox reaction has no effect on the concentrations, so the bulk values are recovered:

$$\begin{aligned} c_A(t, x) &\rightarrow c_A^* \quad \text{as } x \rightarrow \infty, \\ c_B(t, x) &\rightarrow c_B^* \quad \text{as } x \rightarrow \infty. \end{aligned} \tag{2.9}$$

2.2.4 Current

The primary experimental observable of a voltammetry experiment is the current measured by the potentiostat. At any instant the current for the single-electron reduction of species A is

$$I = F A_{\text{el}} j_{0,A}, \tag{2.10}$$

where $A_{\text{el}} = \pi\epsilon^2$ is the area of the disc electrode and $j_{0,A}$ is the net flux of species A at the electrode surface, which is proportional to the concentration gradient at that point by Fick's first law Eq. (2.4). Under our sign convention a positive $j_{0,A}$ corresponds to net consumption of A at the surface (reduction) and hence a positive current. Computing the current therefore reduces to evaluating how $j_{0,A}$ evolves in time.

2.3 Non-Dimensionalisation

The combination of Eqs. (2.3) and (2.6) to (2.9) constitutes an initial-boundary value problem. It is advantageous to solve a non-dimensional version of this problem: the results of a single non-dimensional simulation can be rescaled to any combination of physical parameters without rerunning the computation. We non-dimensionalise in the standard way, scaling concentrations by the bulk concentration of A, diffusion coefficients by D_A , lengths by the electrode radius ϵ , and potentials by the thermal voltage $\mathcal{RT}/F$.

We present a summary of the normalisation in Table 2.1.

Parameter	Normalisation
concentration	$C_j = c_j/c_A^*$
diffusion coefficient	$d_j = D_j/D_A$
spatial coordinate	$X = x/\epsilon$
time	$T = D_A t/\epsilon^2$
applied potential	$\theta = (F/\mathcal{RT}) E$
formal potential	$\theta_f = (F/\mathcal{RT}) E_f^0$
rate constant	$K_0 = k_0 \epsilon/D_A$
scan rate	$\sigma = (\epsilon^2/D_A)(F/\mathcal{RT}) \nu$
current	$J = I/(\pi \epsilon F D_A c_A^*)$

Table 2.1: Dimensionless normalisation of variables. Here ϵ is the radius of the disc electrode.

We define the dimensionless concentration, diffusion coefficient, spatial coordinate, and time as

$$C_j = \frac{c_j}{c_A^*}, \quad d_j = \frac{D_j}{D_A}, \quad X = \frac{x}{\epsilon}, \quad T = \frac{D_A t}{\epsilon^2}, \quad (2.11)$$

and the dimensionless rate constant as

$$K_0 = \frac{k_0 \epsilon}{D_A}. \quad (2.12)$$

We now substitute these into the governing equations, initial conditions, and boundary conditions.

2.3.1 Governing Equation

Considering the left-hand side of Eq. (2.3):

$$\frac{\partial c_j}{\partial t} = \frac{\partial c_j}{\partial C_j} \frac{\partial C_j}{\partial t} = c_A^* \frac{\partial C_j}{\partial t} = \frac{c_A^* D_A}{\epsilon^2} \frac{\partial C_j}{\partial T}.$$

Now the right-hand side:

$$D_j \frac{\partial^2 c_j}{\partial x^2} = D_j \frac{c_A^*}{\epsilon^2} \frac{\partial^2 C_j}{\partial X^2} = d_j \frac{c_A^* D_A}{\epsilon^2} \frac{\partial^2 C_j}{\partial X^2}.$$

Equating gives the dimensionless diffusion equation:

$$\frac{\partial C_j}{\partial T} = d_j \frac{\partial^2 C_j}{\partial X^2}. \quad (2.13)$$

2.3.2 Initial and Far-Field Conditions

The initial conditions become

$$\begin{aligned} C_A(0, X) &= 1, \\ C_B(0, X) &= C_B^*. \end{aligned} \quad (2.14)$$

In practice we cannot simulate a semi-infinite domain. Einstein's analysis of Brownian motion [7] shows that the root-mean-square displacement of a diffusing particle is $x_{\text{rms}} = \sqrt{2Dt}$. To ensure the outer boundary is sufficiently remote that it is unaffected by diffusion, we truncate the domain at $X_{\text{max}} = 6\sqrt{T_{\text{sim}}}$ in dimensionless form. The far-field boundary conditions are then

$$\begin{aligned} C_A(T, X_{\text{max}}) &= 1, \\ C_B(T, X_{\text{max}}) &= C_B^*. \end{aligned} \tag{2.15}$$

2.3.3 Electrode Surface Boundary

We define the dimensionless applied and formal potentials as

$$\begin{aligned} \theta(T) &= \frac{F}{\mathcal{RT}} E(t), \\ \theta_f &= \frac{F}{\mathcal{RT}} E_f^0. \end{aligned} \tag{2.16}$$

Substituting the dimensionless variables into Eqs. (2.6) and (2.7) and using the definition of K_0 from Eq. (2.12), the Butler–Volmer boundary condition and the mass conservation condition become

$$\begin{aligned} \left. \frac{\partial C_A}{\partial X} \right|_{X=0} &= K_0 [C_A(T, 0) e^{-\alpha(\theta(T)-\theta_f)} - C_B(T, 0) e^{(1-\alpha)(\theta(T)-\theta_f)}], \\ d_B \left. \frac{\partial C_B}{\partial X} \right|_{X=0} &= - \left. \frac{\partial C_A}{\partial X} \right|_{X=0}. \end{aligned} \tag{2.17}$$

2.3.4 Scan Rate

As defined in Eq. (2.2), the scan rate is the rate of change of potential with respect to time. We therefore define a dimensionless scan rate σ as the rate of change of dimensionless potential with respect to dimensionless time:

$$\sigma = \frac{d\theta}{dT}. \tag{2.18}$$

To relate σ to ν we note that

$$\nu = \frac{\partial E}{\partial t} = \frac{\mathcal{RT}}{F} \frac{\partial \theta}{\partial t} = \frac{\mathcal{RT}D_A}{F\epsilon^2} \frac{\partial \theta}{\partial T},$$

and hence

$$\sigma = \frac{F\epsilon^2}{\mathcal{RT}D_A} \nu. \tag{2.19}$$

2.3.5 Current

Since the model is now in dimensionless units we cannot compute the flux $j_{0,A}$ directly, but we can compute a dimensionless flux and relate the two. We define

$$J_{X^*,j} = -d_j \left(\frac{\partial C_j}{\partial X} \right)_{X^*}. \quad (2.20)$$

To recover the dimensional flux, we note that

$$j_j = -D_j \left(\frac{\partial c_j}{\partial x} \right)_{x^*} = -d_j \frac{c_A^* D_A}{\epsilon} \frac{\partial C_j}{\partial X} = \frac{c_A^* D_A}{\epsilon} J_j.$$

Substituting into Eq. (2.10) with $A_{\text{el}} = \pi \epsilon^2$ gives

$$I = \pi \epsilon F D_A c_A^* J_{0,A}. \quad (2.21)$$

When we run a simulation we record the dimensionless flux J as a function of the dimensionless potential θ . To obtain the physical voltammogram the results need only be scaled by the values of c_A^* , D_A , and ϵ . Consequently, the results of a single simulation apply to any set of these parameters; the simulation does not need to be rerun when, for example, the bulk concentration of A is changed.

2.4 Equal Diffusion Coefficients

For the remainder of this dissertation we assume that all diffusion coefficients are equal, i.e. $d_B = 1$. This assumption allows us to solve for the concentration of A alone and infer the concentration of B.

From the mass conservation condition Eq. (2.17), when $d_B = 1$ the surface flux condition simplifies to

$$\left. \frac{\partial C_B}{\partial X} \right|_{X=0} = - \left. \frac{\partial C_A}{\partial X} \right|_{X=0}. \quad (2.22)$$

Since the diffusion coefficients are equal and both species satisfy the same diffusion equation, the sum $C_A + C_B$ satisfies

$$\frac{\partial(C_A + C_B)}{\partial T} = \frac{\partial^2(C_A + C_B)}{\partial X^2} \quad (2.23)$$

with homogeneous Neumann conditions at $X = 0$ (by Eq. (2.22)) and constant far-field values. Together with the initial conditions, this implies that $C_A + C_B$ is constant everywhere in space and time. Any decrease in C_A is therefore matched by an equal increase in C_B , and vice versa, so that

$$C_A + C_B = 1 + C_B^*. \quad (2.24)$$

We henceforth drop the subscript and write $C \equiv C_A$, with $C_B = 1 + C_B^* - C$. The electrode boundary condition then reduces to

$$\left. \frac{\partial C}{\partial X} \right|_{X=0} = K_0 [C(T, 0) e^{-\alpha(\theta(T) - \theta_f)} - (1 + C_B^* - C(T, 0)) e^{(1-\alpha)(\theta(T) - \theta_f)}]. \quad (2.25)$$

For ease of notation we denote:

$$\begin{aligned} K_{\text{red}} &= K_0 \exp \left[-\alpha(\theta(T) - \theta_f) \right], \\ K_{\text{ox}} &= K_0 \exp \left[(1 - \alpha)(\theta(T) - \theta_f) \right]. \end{aligned} \quad (2.26)$$

These shorthand expressions are used extensively in the numerical discretisation that follows in next chapter.

Chapter 3

Numerical Methods

In this chapter we discretise the dimensionless initial-boundary value problem derived in the previous chapter and describe how the resulting linear system is solved at each timestep. The reader is assumed to be familiar with finite difference methods; we focus on the aspects specific to the electrochemical setting, namely the exponentially expanding spatial mesh and the incorporation of the Butler–Volmer boundary condition into the tridiagonal system.

For the remainder of this chapter we assume $C_B^* = 0$ (no product initially present in solution), so that the equal-diffusion-coefficient boundary condition Eq. (2.25) becomes

$$\left. \frac{\partial C}{\partial X} \right|_{X=0} = K_{\text{red}}(T) C(T, 0) - K_{\text{ox}}(T) (1 - C(T, 0)), \quad (3.1)$$

where K_{red} and K_{ox} were defined in Eq. (2.26) and depend on time through the applied potential $\theta(T)$.

3.1 Discretisation

3.1.1 Temporal Grid

We divide the simulation interval $[0, T_{\text{max}}]$ into m uniform steps of size

$$\Delta T = \frac{T_{\text{max}}}{m}, \quad (3.2)$$

so that the discrete time levels are $T^k = k \Delta T$ for $k = 0, 1, \dots, m$. Here T_{max} is determined by the scan rate and the potential window of the experiment: for a cyclic voltammogram sweeping from θ_i to θ_v and back, $T_{\text{max}} = \frac{2|\theta_v - \theta_i|}{\sigma}$.

3.1.2 Spatial Grid

The electrochemical reaction occurs at the electrode surface and its influence on the concentration profile is confined to a thin diffusion layer near $X = 0$. Far from the electrode the concentration remains at its bulk value and contributes little to the dynamics. It is therefore desirable to concentrate spatial resolution near the electrode surface whilst using coarser spacing in the far field. We achieve this with an exponentially expanding mesh controlled by two parameters: the initial spacing $h_0 > 0$ and an expansion factor $\omega > 1$.

We define n spatial nodes $\{X_0, X_1, \dots, X_{n-1}\}$ by

$$X_i = \begin{cases} \sum_{k=0}^{i-1} h_0 \omega^k, & 0 \leq i \leq n-2, \\ X_{\max}, & i = n-1, \end{cases} \quad (3.3)$$

with $X_0 = 0$ corresponding to the electrode surface. The local spacings about point i are

$$\Delta X_+ = X_{i+1} - X_i, \quad \Delta X_- = X_i - X_{i-1}. \quad (3.4)$$

For $0 \leq i \leq n-2$ the spacing grows geometrically as $\Delta X_+ = h_0 \omega^i$, so the mesh is finest at the electrode and coarsens exponentially into the bulk. The final point ($i = n-1$) is placed at $X_{\max} = 6\sqrt{T_{\max}}$ as discussed in the previous chapter, ensuring the far-field boundary does not influence the solution.

3.2 Finite Difference Scheme

We discretise the dimensionless diffusion equation Eq. (2.13) (with $d_j = 1$) using a backward Euler (fully implicit) scheme in time and a three-point central difference approximation for the second spatial derivative on the non-uniform mesh.

Writing $C_i^k \approx C(T^k, X_i)$, the scheme at an interior node $1 \leq i \leq n-2$ reads

$$\frac{C_i^k - C_i^{k-1}}{\Delta T} = \frac{1}{\frac{1}{2}(\Delta X_+ + \Delta X_-)} \left(\frac{C_{i+1}^k - C_i^k}{\Delta X_+} - \frac{C_i^k - C_{i-1}^k}{\Delta X_-} \right). \quad (3.5)$$

The implicit treatment evaluates the spatial operator at the new time level k , which makes the scheme unconditionally stable and avoids a restrictive CFL condition on the timestep.

3.2.1 Interior Coefficients

Rearranging Eq. (3.5) so that all unknowns at time level k appear on the left and the known value at level $k - 1$ appears on the right, we obtain

$$\alpha_i C_{i-1}^k + \beta_i C_i^k + \gamma_i C_{i+1}^k = C_i^{k-1}, \quad 1 \leq i \leq n - 2, \quad (3.6)$$

where the tridiagonal coefficients are

$$\begin{aligned} \alpha_i &= -\frac{2 \Delta T}{\Delta X_- (\Delta X_+ + \Delta X_-)}, \\ \gamma_i &= -\frac{2 \Delta T}{\Delta X_+ (\Delta X_+ + \Delta X_-)}, \\ \beta_i &= 1 - \alpha_i - \gamma_i. \end{aligned} \quad (3.7)$$

3.2.2 Electrode Boundary ($i = 0$)

At the electrode surface we approximate the spatial gradient in Eq. (3.1) with a two-point forward difference:

$$\frac{C_1^k - C_0^k}{h_0} = K_{\text{red}}^k C_0^k - K_{\text{ox}}^k (1 - C_0^k), \quad (3.8)$$

where K_{red}^k and K_{ox}^k denote the values of K_{red} and K_{ox} evaluated at the potential $\theta(T^k)$. Rearranging into the form $\beta_0 C_0^k + \gamma_0 C_1^k = b_0$ gives

$$\begin{aligned} \beta_0 &= 1 + h_0 K_{\text{red}}^k + h_0 K_{\text{ox}}^k, \\ \gamma_0 &= -1, \\ b_0 &= h_0 K_{\text{ox}}^k. \end{aligned} \quad (3.9)$$

3.2.3 Far-Field Boundary ($i = n - 1$)

The Dirichlet condition $C(T, X_{\text{max}}) = 1$ is imposed directly:

$$\begin{aligned} \beta_{n-1} &= 1, \\ b_{n-1} &= 1, \end{aligned} \quad (3.10)$$

with $\alpha_{n-1} = \gamma_{n-1} = 0$.

3.2.4 Solution of the Tridiagonal System

At each time level k the discretisation produces a tridiagonal linear system with lower diagonal $(\alpha_1, \dots, \alpha_{n-1})$, main diagonal $(\beta_0, \dots, \beta_{n-1})$, and upper diagonal $(\gamma_0, \dots, \gamma_{n-2})$.

Because the electrode boundary coefficients depend on the applied potential, the first row of the system is reassembled at each timestep; the interior coefficients are time-independent and can be precomputed.

This system is solved in $O(n)$ operations using the Thomas algorithm (tridiagonal LU factorisation), making the cost per timestep linear in the number of spatial nodes.

3.3 The Forward Solver as a Mapping

The implicit finite difference solver computes the concentration field at every node in the spatial-temporal mesh. However, the experimentally observable quantity is not the concentration profile but the current, which is proportional to the surface flux $J_{0,A}$ recorded at each timestep. It is therefore natural to regard the solver as a mapping from the set of model parameters to the predicted dimensionless current response.

We define the *forward map*

$$\mathcal{P}: \Phi \rightarrow \mathcal{J}, \quad \phi \mapsto \mathcal{P}(\phi), \quad (3.11)$$

where $\Phi \subseteq \mathbb{R}^3$ is the parameter space with $\phi = (\alpha, K_0, \theta_f)$, and $\mathcal{J}$ is the space of dimensionless flux trajectories $J_{0,A}(T)$ evaluated at the discrete time levels. Given a parameter vector ϕ , a single evaluation of $\mathcal{P}$ consists of assembling and solving the tridiagonal system at each of the m timesteps and extracting the surface flux at each step.

Although the electrode boundary condition is discretised with a two-point approximation (to preserve the tridiagonal structure), the surface flux used to compute the current is evaluated with a higher-order three-point formula. Writing $h_1 = X_1 - X_0$ and $h_2 = X_2 - X_0$, we fit a quadratic through (X_0, C_0^k) , (X_1, C_1^k) , (X_2, C_2^k) and differentiate at X_0 to obtain

$$J_{0,A}^k = - \left(\frac{\partial C}{\partial X} \right) \Big|_{X=0} \approx - \frac{h_2^2 (C_1^k - C_0^k) - h_1^2 (C_2^k - C_0^k)}{h_1 h_2 (h_2 - h_1)}. \quad (3.12)$$

This forward map is the central object in the inference problems discussed in the next chapter: given experimental data, we seek the parameters ϕ for which $\mathcal{P}(\phi)$ best reproduces the observed current response.

3.4 Effect of Butler–Volmer Parameters

In the electron-transfer reaction we have been discussing there are three parameters of interest: the transfer coefficient α , the dimensionless rate constant K_0 , and the

dimensionless formal potential θ_f . Each parameter has a distinct qualitative effect on the shape of the voltammogram. We illustrate these effects by varying each parameter in turn about a quasi-reversible baseline in Fig. 3.1.

Increasing the transfer coefficient α enhances the cathodic peak current whilst diminishing the anodic peak; decreasing α has the opposite effect. This asymmetry arises because α weights the reduction and oxidation exponentials in the Butler–Volmer equation: a larger α favours the cathodic branch at the expense of the anodic branch. The dimensionless rate constant K_0 has a symmetrising effect on the voltammogram. As K_0 increases, the cathodic and anodic peaks sharpen, grow in magnitude, and move closer together, reflecting the transition from kinetically limited to near-reversible behaviour. Finally, the formal potential θ_f translates the entire voltammogram along the potential axis without altering the peak shapes, since it simply shifts the point about which the overpotential is measured.




Figure 3.1: Effect of each Butler–Volmer parameter on the dimensionless voltammogram for a quasi-reversible reaction. In each panel one parameter is varied while the others are held fixed. Left: varying α changes the relative heights of the cathodic and anodic peaks. Centre: increasing K_0 symmetrises and sharpens the voltammogram. Right: varying θ_f translates the peaks along the potential axis.

The interplay between these parameters is important to note. Because changes in one parameter can produce features in the voltammogram that partially resemble changes in another, there is potential for compensating errors in parameter estimation. For instance, a rate constant that is too small — which broadens the peaks and shifts them apart — could be partially offset by an adjustment in the formal potential that translates the peaks back towards the observed positions. Such parameter correlations suggest that a point estimate obtained by optimisation alone may not capture the full uncertainty in the inferred values.

The situation is compounded in the reversible limit. Figure 3.2 shows the effect of varying α at $K_0 = 200$, well into the reversible regime. The three voltammograms for $\alpha = 0.3, 0.5, 0.7$ are nearly indistinguishable: when the rate constant is large the exponential terms in the Butler–Volmer equation force the surface concentrations to their thermodynamic equilibrium values regardless of α , and the transfer coefficient effectively drops out of the problem. In this regime, any attempt to estimate α from the voltammogram alone — whether by maximum likelihood or other point-estimation methods — will yield an essentially uninformative result.




Figure 3.2: Effect of the transfer coefficient α on the voltammogram in the reversible regime ($K_0 = 200$). The three curves are nearly indistinguishable, demonstrating that α has negligible influence on the voltammogram when the reaction kinetics are fast.

The combination of parameter correlations in the quasi-reversible regime, the practical non-identifiability of α in the reversible regime, and the inherent noise present in experimental data means that estimating the Butler–Volmer parameters is a problem better suited to a probabilistic framework than to a purely optimisation-based approach. A probabilistic treatment naturally quantifies the uncertainty in each parameter and reveals correlations between them, providing a more complete picture than a single point estimate. We develop this approach in the following chapter.

Chapter 4

Bayesian Approach to the Inverse Problem

In general, data obtained from a voltammetry experiment will be noisy and subject to experimental error. Further, as discussed in Section 3.4, there may be multiple configurations of parameters that produce similar current waveforms, and certain parameters may have negligible influence on the voltammogram in particular kinetic regimes. This makes the problem of recovering the parameters of the experiment given the observed data – the *inverse problem* – *ill-posed*: a unique, stable solution may not exist. For this reason, we take a Bayesian approach and seek probability distributions over the parameters rather than point estimates. This framework naturally quantifies the uncertainty in each parameter and reveals correlations between them.

4.1 Noise Model

We model the observed dimensionless flux data as the output of the forward map $\mathcal{P}$ corrupted by additive Gaussian noise. Let $\mathbf{y} = (y_1, \dots, y_N)$ denote the N flux measurements recorded at discrete time levels. We assume

$$y_i = \mathcal{P}(\phi)_i + \varepsilon_i, \quad \varepsilon_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \varsigma^2), \quad (4.1)$$

where $\mathcal{P}(\phi)_i$ is the i th component of the simulated flux trajectory for parameter vector $\phi = (\alpha, K_0, \theta_f)$, and ς^2 is the variance of the measurement noise. We use ς rather than σ to avoid confusion with the dimensionless scan rate defined in the previous chapters.

We set standard deviation of the measurement noise is set as a fraction of the peak flux magnitude. Let ϕ^* denote the true parameter vector used to generate the

synthetic data. We define

$$\varsigma = \eta |J_{\text{peak}}|, \quad J_{\text{peak}} = \min_i \mathcal{P}(\phi^*)_i, \quad (4.2)$$

where $\eta \in (0, 1)$ is a dimensionless noise parameter. Experimental measurements typically exhibit noise levels of 1–2% of the peak current [18], corresponding to $\eta \in [0.01, 0.02]$. Since the cathodic flux is negative, the peak magnitude is given by the minimum of the noiseless simulated trajectory. For a given η , the noise level ς is determined entirely by the noiseless forward solve and is treated as a known constant throughout the inference.

An example of such noisy data is shown in Fig. 4.1.




Figure 4.1: Synthetic voltammogram data generated by evaluating the forward map $\mathcal{P}$ with dimensionless scan rate $\sigma = 1000$ over the potential window $\theta_i = 10$ to $\theta_v = -10$, then adding i.i.d. Gaussian noise. The noisy observations are the data $\mathbf{y}$ from which the Butler–Volmer parameters are to be inferred.

4.2 Posterior Distribution

Bayes’ theorem relates the posterior distribution of the parameters given the data to the likelihood and prior:

$$p(\phi \mid \mathbf{y}) = \frac{L(\mathbf{y} \mid \phi) \pi(\phi)}{p(\mathbf{y})}, \quad (4.3)$$

where $L(\mathbf{y} \mid \phi)$ is the likelihood, $\pi(\phi)$ is the prior distribution, and $p(\mathbf{y}) = \int L(\mathbf{y} \mid \phi) \pi(\phi) d\phi$ is the marginal likelihood (or evidence), which serves as a normalising constant independent of ϕ .

4.2.1 Prior

We adopt an uninformative prior: a uniform distribution over physically reasonable bounds for each parameter,

$$\pi(\phi) \propto 1 \quad \text{for } \phi \in \Phi, \quad (4.4)$$

where $\Phi \subset \mathbb{R}^3$ is a bounded region encoding constraints such as $\alpha \in (0, 1)$ and $K_0 > 0$. Within Φ the prior is constant and therefore does not favour any particular parameter combination. Consequently, the posterior is proportional to the likelihood:

$$p(\phi \mid \mathbf{y}) \propto L(\mathbf{y} \mid \phi). \quad (4.5)$$

4.2.2 Likelihood

Under the noise model Eq. (4.1), the observations are conditionally independent given ϕ , and the likelihood takes the form

$$L(\mathbf{y} \mid \phi) = \prod_{i=1}^N \frac{1}{\sqrt{2\pi\varsigma^2}} \exp\left(-\frac{(y_i - \mathcal{P}(\phi)_i)^2}{2\varsigma^2}\right). \quad (4.6)$$

In practice we work with the log-likelihood. Taking the logarithm of Eq. (4.6) gives

$$\log L(\mathbf{y} \mid \phi) = -\frac{N}{2} \log(2\pi\varsigma^2) - \frac{1}{2\varsigma^2} \sum_{i=1}^N (y_i - \mathcal{P}(\phi)_i)^2. \quad (4.7)$$

Since ς is treated as known, the first term is a constant that does not depend on ϕ and can be dropped when comparing likelihoods at different parameter values. We therefore define the log-likelihood function used in all subsequent computations as

$$\ell(\phi) = -\frac{1}{2\varsigma^2} \sum_{i=1}^N (y_i - \mathcal{P}(\phi)_i)^2. \quad (4.8)$$

4.3 Markov Chain Monte Carlo

The posterior distribution Eq. (4.5) is known only up to a normalising constant, and the integral defining that constant is generally intractable for non-trivial forward maps. *Markov chain Monte Carlo* (MCMC) provides a numerical approach: rather than evaluating the posterior everywhere, we construct a Markov chain whose stationary distribution is the posterior and use the samples generated by the chain to approximate expectations and marginal distributions.

4.3.1 Markov Chain

An MCMC method produces a sequence of parameter vectors

$$\{\phi^{(0)}, \phi^{(1)}, \dots, \phi^{(T)}\} \subseteq \Phi, \quad (4.9)$$

where each $\phi^{(t+1)}$ depends only on the current state $\phi^{(t)}$ (the Markov property). Under regularity conditions — in particular, that the chain is irreducible and aperiodic — the ergodic theorem guarantees that the empirical distribution of the samples converges to the target posterior as $T \rightarrow \infty$.

4.3.2 Proposal Distribution

At each step the chain generates a *candidate* parameter vector from a *proposal distribution* $q(\cdot \mid \phi^{(t)})$:

$$\phi_{\text{cand}} \sim q(\cdot \mid \phi^{(t)}). \quad (4.10)$$

The proposal distribution is a design choice that determines how the chain explores the parameter space. Its form need not match the posterior, but the efficiency of the sampler depends heavily on how well the proposals reflect the local geometry of the target distribution.

4.3.3 Acceptance Criterion

Whether the chain moves to the candidate is determined by the *Metropolis–Hastings acceptance probability* [13]:

$$r = \min \left\{ \frac{p(\phi_{\text{cand}} \mid \mathbf{y}) q(\phi^{(t)} \mid \phi_{\text{cand}})}{p(\phi^{(t)} \mid \mathbf{y}) q(\phi_{\text{cand}} \mid \phi^{(t)})}, 1 \right\}. \quad (4.11)$$

A uniform random number $u \sim \text{Uniform}(0, 1)$ is drawn: if $u < r$ the candidate is accepted and $\phi^{(t+1)} = \phi_{\text{cand}}$; otherwise the chain remains at $\phi^{(t+1)} = \phi^{(t)}$. The ratio in Eq. (4.11) involves the posterior only through ratios, so the unknown normalising constant cancels – this is the key property that makes MCMC practical.

4.4 Random Walk Metropolis–Hastings

The *random walk Metropolis–Hastings* (RWMH) algorithm is a particular instance of the Metropolis–Hastings framework in which the proposal distribution is a symmetric Gaussian centred at the current state:

$$q(\phi_{\text{cand}} \mid \phi^{(t)}) = \mathcal{N}(\phi^{(t)}, \Sigma), \quad (4.12)$$

where $\Sigma \in \mathbb{R}^{3 \times 3}$ is a positive-definite covariance matrix that controls the size and orientation of the proposed steps. Since the Gaussian density is symmetric — $q(\phi_{\text{cand}} | \phi^{(t)}) = q(\phi^{(t)} | \phi_{\text{cand}})$ — the proposal ratio in Eq. (4.11) is unity, and the acceptance probability simplifies to

$$r = \min \left\{ \frac{p(\phi_{\text{cand}} | \mathbf{y})}{p(\phi^{(t)} | \mathbf{y})}, 1 \right\} = \min \{ \exp(\ell(\phi_{\text{cand}}) - \ell(\phi^{(t)})), 1 \}, \quad (4.13)$$

where the second equality uses Eq. (4.5) and the definition of the log-likelihood Eq. (4.8). In practice, we compute $\ell(\phi_{\text{cand}}) - \ell(\phi^{(t)})$ directly to avoid numerical overflow.

The choice of the proposal covariance Σ is critical to the efficiency of the sampler. If the steps are too small the chain explores the parameter space slowly and successive samples are highly correlated; if the steps are too large most proposals fall in regions of low posterior density and are rejected. A common heuristic is to tune Σ to achieve an acceptance rate of approximately 23% [10], which is asymptotically optimal for Gaussian targets in high dimensions.

4.4.1 Sampling Results

We apply the RWMH algorithm to synthetic data generated from the forward map with known true parameters and additive Gaussian noise with $\varsigma = 0.25$. Figure 4.2 shows the marginal posterior histograms for each parameter. All three marginals concentrate around the true parameter values, confirming that the sampler recovers the correct region of the parameter space. The marginal for K_0 and θ_f are sharply peaked, indicating that the data is informative for these parameters; the marginal for α is broader, consistent with the weaker sensitivity discussed in Section 3.4.

However, the histograms also exhibit extended tails, particularly for α . The origin of these tails is visible in Fig. 4.3, which shows the (α, K_0) projection of the chain coloured by sample index. The chain was initialised far from the high-density region and spends a substantial number of iterations traversing a low-probability ridge before settling near the true values. The standard practice in MCMC is to discard these initial samples as *burn-in* and retain only the remainder as approximate draws from the posterior. In principle this is straightforward; in our setting, however, each likelihood evaluation requires a full run of the forward map $\mathcal{P}$, which involves solving the tridiagonal system at every timestep. The cost of generating each proposal is therefore non-trivial, and discarding a large burn-in period represents a significant waste of computational effort.




Figure 4.2: Marginal posterior histograms for the three Butler–Volmer parameters obtained by RWMH sampling on synthetic data with noise level $\varsigma = 0.25$. Dashed lines indicate the true parameter values used to generate the data. The posteriors for K_0 and θ_f are tightly concentrated around the true values, while the marginal for α is broader, reflecting its weaker influence on the voltammogram at this kinetic regime.

A more efficient strategy would be to first locate the high-density region of the posterior — effectively maximising the likelihood — and then initialise the Markov chain there, so that all samples contribute to the posterior approximation from the outset. Gradient-based optimisation methods are well suited to this task, as they exploit the local slope of the likelihood surface to move rapidly toward the mode. More ambitiously, gradient information can be incorporated into the MCMC proposal itself, yielding samplers that explore the posterior far more efficiently than the random walk. Both of these ideas require the ability to differentiate the log-likelihood with respect to the parameters, which in turn requires differentiating through the forward map $\mathcal{P}$. We develop this capability in the following chapter.




Figure 4.3: Scatter plot of the RWMH chain projected onto the (α, K_0) plane. Points are coloured by sample index: early samples (dark) show the chain traversing a low-density region far from the true parameters (black cross), while later samples (light) concentrate near the posterior mode. The extended transient illustrates the cost of burn-in when each likelihood evaluation requires a full forward solve.

Chapter 5

Differentiating Through the Forward Solver

The previous chapter postulated that gradient information is could be desirable both for initialising the Markov chain and for constructing more efficient MCMC proposals. Computing the gradient of the log-likelihood Eq. (4.8) with respect to the parameters ϕ requires differentiating through the forward map $\mathcal{P}$, which itself involves solving a sequence of tridiagonal linear systems. In this chapter we describe how these derivatives are obtained efficiently.

5.1 Discretise Then Optimise

There are broadly two strategies for differentiating through a numerical solver. In the *optimise then discretise* approach one derives the continuous adjoint equations analytically and then discretises them; in the *discretise then optimise* approach one first discretises the forward problem and then differentiates the discrete computation. Generally, the discretise then optimise is preferable as it computes the exact gradient of the quantity that is actually evaluated in code, avoiding the consistency errors that can arise when the adjoint PDE is discretised independently of the forward PDE [16]. We adopt this approach throughout.

Concretely, our forward map $\mathcal{P}$ is implemented as a sequence of algebraic operations — assembling tridiagonal systems and solving them via the Thomas algorithm at each timestep, then extracting the surface flux. The gradient we seek is therefore the derivative of a composition of these discrete operations with respect to the parameters ϕ .

5.2 Reverse Mode Automatic Differentiation

The log-likelihood is a scalar-valued function of the parameter vector. In this setting reverse mode automatic differentiation (AD) is the natural choice, as it computes the full gradient in a single backward pass regardless of the number of parameters.

Reverse mode AD decomposes the gradient computation into a sequence of *vector-Jacobian products* (VJPs). If the forward computation is written as a composition $f = f_L \circ f_{L-1} \circ \dots \circ f_1$, then the chain rule gives

$$\frac{\partial f}{\partial \phi} = \frac{\partial f_L}{\partial f_{L-1}} \cdot \frac{\partial f_{L-1}}{\partial f_{L-2}} \dots \frac{\partial f_1}{\partial \phi}. \quad (5.1)$$

In reverse mode, this product is evaluated from left to right: at each stage a row vector (the adjoint, or cotangent) is multiplied by the Jacobian of the corresponding primitive operation. The intermediate Jacobians are never formed explicitly; only the VJP of each primitive is required.

For the forward map $\mathcal{P}$ the dominant primitives are the tridiagonal solves — one per timestep. The question therefore reduces to: how should the VJP of a tridiagonal solve be implemented?

5.3 Differentiating Through Linear Solves

Consider a single linear system $Ax = b$, where A depends on the parameters ϕ and b depends on the solution at the previous timestep. A naïve approach is to record every intermediate value of the Thomas algorithm on the forward pass and replay them in reverse during the backward pass. This is what a generic AD framework will do if the Thomas algorithm is implemented as a sequence of elementary operations: the entire computational graph of the factorisation and back-substitution is stored on a tape.

For a single tridiagonal solve on n unknowns the memory overhead of taping is modest. However, the forward map involves m consecutive solves (one per timestep), and the tape must be retained across all of them because the backward pass traverses the full time history in reverse. The total memory cost therefore scales as $O(mn)$, which becomes prohibitive for fine spatial and temporal grids. Worse, the overhead of maintaining and traversing the tape degrades runtime performance even when memory is not the binding constraint.

5.4 Adjoint of the Tridiagonal Solve

A more efficient approach exploits the structure of the linear solve directly. Suppose the forward pass solves $Ax = b$ and the backward pass receives the adjoint (cotangent) $\bar{x} = \partial\ell/\partial x$. We require the adjoint with respect to b :

$$\bar{b} = \frac{\partial\ell}{\partial b} = \bar{x}^T \frac{\partial x}{\partial b} = \bar{x}^T A^{-1}. \quad (5.2)$$

This is equivalent to solving the *adjoint system*

$$A^T \lambda = \bar{x}, \quad (5.3)$$

and setting $\bar{b} = \lambda^T$. Similarly, the adjoint with respect to the matrix entries of A is

$$\bar{A} = -\lambda x^T, \quad (5.4)$$

which follows from differentiating the identity $Ax = b$ and substituting the solution of the adjoint system.

The key observation is that A^T is also tridiagonal (with the sub- and super-diagonals transposed), so the adjoint system Eq. (5.3) can be solved by the Thomas algorithm at the same $O(n)$ cost as the forward solve. Crucially, this approach requires only the matrix A and the forward solution x to be stored — not the full tape of intermediate values from the factorisation. The memory cost per timestep is therefore $O(n)$ rather than the $O(n)$ per operation of naïve taping, and the total cost across all m timesteps is reduced from taping the entire forward history to storing only the tridiagonal coefficients and solutions at each level.

In practice we implement this by defining a custom VJP rule for the tridiagonal solve in *JAX* [2]. On the forward pass the Thomas algorithm solves $Ax = b$ and caches A and x . On the backward pass the custom rule receives $\bar{x}$, solves $A^T \lambda = \bar{x}$ using the banded transpose, and returns $\bar{b} = \lambda^T$ and $\bar{A} = -\lambda x^T$. No taping of the Thomas algorithm’s internal operations is required.

5.5 Extension to Pentadiagonal Systems

The benefits of the custom adjoint become even more pronounced for higher-order discretisations that produce pentadiagonal (or wider banded) systems. A pentadiagonal factorisation involves substantially more intermediate values than a tridiagonal one, and taping the full forward computation through the banded solver would incur a correspondingly larger memory and runtime penalty. The adjoint approach sidesteps

this entirely: the transpose of a pentadiagonal matrix is again pentadiagonal, so the adjoint system is solved by a banded solver at the same $O(n)$ cost per timestep. The custom VJP rule generalises directly — one simply solves with the transposed band structure — making this a crucial step in enabling the use of higher-order spatial discretisations within the differentiable forward map.

5.6 Initialisation Results

As discussed in Section 4.4.1, initialising a Markov chain far from the posterior mode wastes computational resources on burn-in samples that must be discarded. A gradient-based optimisation stage can locate the high-density region before sampling begins, so that the chain starts collecting useful draws immediately. A good initialisation strategy should satisfy three criteria: it should be robust to the choice of starting point, it should perform reliably across different noise levels in the data, and it should remain stable once it has reached a region of high likelihood rather than oscillating away from it.

We compare the ADAM optimiser [17], which exploits the gradient of the log-likelihood computed via the adjoint method, against CMA-ES [12], a gradient-free evolutionary strategy. CMA-ES maintains a population of candidate solutions; we set the population size to 4. Since the backward (adjoint) pass costs approximately $1\text{--}3\times$ the forward pass, a single ADAM step (one forward and one backward pass) is roughly comparable in cost to evaluating the forward map for all four members of the CMA-ES population. We therefore equalise the computational budget by giving both algorithms the same number of iterations, each iteration comprising one ADAM update or one CMA-ES generation of four candidates.

Figure 5.1 shows the log-posterior density over the course of the optimisation at two noise levels ($\eta = 0.01$ and $\eta = 0.02$). Each algorithm is run from 32 initial conditions drawn uniformly from a physically reasonable region of the parameter space; solid lines show the mean and shaded regions indicate one standard deviation across these runs. Both methods converge toward the posterior mode, but ADAM does so more rapidly and with markedly lower variance across initialisations, indicating greater robustness to starting conditions.




Figure 5.1: Log-posterior density during optimisation for ADAM (blue) and CMA-ES at noise levels $\eta = 0.02$ and $\eta = 0.01$. Each algorithm is run from 32 randomly drawn initial conditions; solid lines show the mean and shaded regions indicate ± 1 standard deviation across runs. The dashed red line marks the log-density evaluated at five times the true parameters – denoting a high density area. Both algorithms are given equal computational budgets (CMA-ES population size 4, matching the cost of one ADAM gradient step). ADAM converges faster, with lower variance across initialisations, and remains stable once it reaches the high-density region.

Chapter 6

Hamiltonian Monte Carlo

The random walk Metropolis–Hastings algorithm explored in Chapter 4 generates proposals by perturbing the current state with isotropic Gaussian noise. This strategy ignores the geometry of the posterior: proposals are equally likely in every direction, regardless of whether the posterior is steep or flat. Now that we have the ability to differentiate through the forward map, we can exploit gradient information to guide the sampler along the posterior surface. *Hamiltonian Monte Carlo* (HMC) [1, 19] does precisely this by borrowing the formalism of Hamiltonian dynamics to construct proposals that are both distant from the current state and likely to be accepted.

The physical intuition is as follows. We imagine the current position of the Markov chain as a particle moving in the parameter space. At each step, the particle is given a random kick (a momentum draw) and then allowed to evolve under a Hamiltonian that couples its position to the posterior. Because the dynamics conserve energy, the particle naturally spends more time in regions of high posterior density and traces out long-range, low-energy trajectories that would take a random walk many steps to traverse.

6.1 The Hamiltonian

Let $q \in \mathbb{R}^d$ denote the position (i.e. the parameter vector ϕ ; we use q to follow the Hamiltonian mechanics convention) and introduce an auxiliary momentum variable $p \in \mathbb{R}^d$. We define the joint density of position and momentum through the canonical distribution

$$P(q, p) \propto \exp(-H(q, p)), \tag{6.1}$$

where $H(q, p)$ is the Hamiltonian. We decompose H into a potential energy $U(q)$ and a kinetic energy $K(p)$:

$$H(q, p) = U(q) + K(p). \quad (6.2)$$

This separation ensures that the joint density factorises as $P(q, p) \propto \exp(-U(q)) \exp(-K(p))$, so the position and momentum are independent: marginalising over p recovers the target distribution over q alone.

6.1.1 Potential Energy

We define the potential energy as the negative log-posterior:

$$U(q) = -\log p(q \mid \mathbf{y}) = -\ell(q) + \text{const}, \quad (6.3)$$

where the constant absorbs the normalising factor. Under this identification, regions of high posterior density correspond to low potential energy. The physical picture is then that of a landscape in which the particle rolls into valleys at the posterior modes and must climb uphill to escape — precisely the behaviour we want from a sampler.

6.1.2 Kinetic Energy

The standard choice of kinetic energy is the quadratic form

$$K(p) = \frac{1}{2} p^T M^{-1} p, \quad (6.4)$$

where $M \in \mathbb{R}^{d \times d}$ is a symmetric positive-definite *mass matrix*. Under Eq. (6.1) the momentum is then distributed as $p \sim \mathcal{N}(0, M)$. The mass matrix controls the inertia of the particle in each direction: a well-chosen M that approximates the posterior covariance can precondition the dynamics so that the particle explores all directions at a similar rate.

6.2 Algorithm

Hamilton’s equations of motion for the system Eq. (6.2) are

$$\frac{dq}{dt} = \frac{\partial H}{\partial p} = M^{-1} p, \quad \frac{dp}{dt} = -\frac{\partial H}{\partial q} = -\nabla_q U(q). \quad (6.5)$$

The gradient $\nabla_q U(q) = -\nabla_q \ell(q)$ is the derivative of the log-likelihood with respect to the parameters, computed via the adjoint method described in Chapter 5.

6.2.1 Leapfrog Integration

In practice Hamilton's equations are integrated numerically using the *leapfrog* integrator, which is symplectic and time-reversible. Given a step size ε and L integration steps, one leapfrog trajectory from (q_0, p_0) proceeds as:

$$\begin{aligned} p_{i+1/2} &= p_i - \frac{\varepsilon}{2} \nabla_q U(q_i), \\ q_{i+1} &= q_i + \varepsilon M^{-1} p_{i+1/2}, \\ p_{i+1} &= p_{i+1/2} - \frac{\varepsilon}{2} \nabla_q U(q_{i+1}), \end{aligned} \tag{6.6}$$

for $i = 0, 1, \dots, L - 1$. Each step requires one evaluation of $\nabla_q U$, so a trajectory of L steps costs L gradient evaluations, each of which involves a full forward and adjoint solve of the model.

6.2.2 Metropolis Correction

Let (q^*, p^*) denote the endpoint of the leapfrog trajectory starting from (q, p) . Because the leapfrog integrator introduces a discretisation error, the Hamiltonian is not exactly conserved along the trajectory. To ensure the chain targets the correct stationary distribution, a Metropolis acceptance step is applied. The sign of the momentum at the endpoint is first negated, $p^* \rightarrow -p^*$, to make the proposal a time-reversal of the forward trajectory; this is necessary to satisfy the detailed balance condition. The acceptance probability is then

$$r = \min\{1, \exp(H(q, p) - H(q^*, p^*))\}. \tag{6.7}$$

Because the leapfrog integrator is symplectic, the energy error is typically small and the acceptance rate is high even for proposals that are far from the current state.

Note that after the accept/reject decision is made the momentum is discarded; the position q (accepted or not) is retained as the next sample of the parameter vector. At the start of the following iteration a fresh momentum is drawn from $\mathcal{N}(0, M)$, which randomises the energy level and ensures the particle explores the full posterior rather than tracing a single energy contour throughout.

The full HMC procedure is summarised in Algorithm 1.

6.3 Sampling Results

Any MCMC sampler produces serially correlated samples: each proposal depends on the current state of the chain, so successive draws are not independent. Correlated

Algorithm 1 Hamiltonian Monte Carlo

Require: Initial state $q^{(0)}$, mass matrix M , step size ε , number of leapfrog steps L , number of iterations T

```
1: for  $t = 0, 1, \dots, T - 1$  do
2:   Draw  $p \sim \mathcal{N}(0, M)$ 
3:   Set  $(q_0, p_0) \leftarrow (q^{(t)}, p)$ 
4:   for  $i = 0, 1, \dots, L - 1$  do
5:     Perform leapfrog step Eq. (6.6):  $(q_i, p_i) \rightarrow (q_{i+1}, p_{i+1})$ 
6:   end for
7:   Set  $(q^*, p^*) \leftarrow (q_L, -p_L)$  ▷ Negate momentum
8:   Compute  $\Delta H = H(q^{(t)}, p) - H(q^*, p^*)$ 
9:   Draw  $u \sim \text{Uniform}(0, 1)$ 
10:  if  $u < \min\{1, \exp(\Delta H)\}$  then
11:     $q^{(t+1)} \leftarrow q^*$ 
12:  else
13:     $q^{(t+1)} \leftarrow q^{(t)}$ 
14:  end if
15: end for
```

Ensure: Chain $\{q^{(0)}, q^{(1)}, \dots, q^{(T)}\}$

samples carry less information about the target distribution than an equal number of independent draws. It is therefore insufficient to report only the raw number of samples; we must also quantify how much independent information the chain has accumulated. The standard diagnostic for this purpose is the *effective sample size* (ESS) [9, 11], defined as

$$\text{ESS} = \frac{T}{1 + 2 \sum_{k=1}^{\infty} \rho_k}, \quad (6.8)$$

where T is the total number of samples and ρ_k is the autocorrelation of the chain at lag k . When consecutive samples are nearly independent ρ_k decays rapidly and $\text{ESS} \approx T$; when the chain is highly correlated the sum in the denominator is large and the effective sample size is much smaller than T .

Figure 6.1 compares the growth of ESS over the sampling duration for HMC and RWMH. Because each HMC sample is more expensive to generate than an RWMH sample — requiring multiple leapfrog steps, each involving a forward and adjoint solve — the comparison is made on a normalised time axis so that the two algorithms are compared at equal computational cost: the HMC chain comprises 32 000 samples and the RWMH chain 160 000 samples. HMC accumulates effective samples at a consistently higher rate across all three parameters, confirming that the gradient information embedded in the Hamiltonian dynamics produces less correlated samples

and a more efficient exploration of the posterior. The advantage is most pronounced for K_0 , where the HMC chain achieves roughly three times the ESS of RWMH at the end of the sampling budget.




Figure 6.1: Effective sample size as a function of normalised computational time for HMC (blue) and RWMH (orange), for each of the three Butler–Volmer parameters. The horizontal axis is normalised so that the two algorithms are compared at equal computational cost (32 000 HMC samples and 160 000 RWMH samples). HMC achieves a higher ESS across all parameters, with the largest gain for K_0 .

6.4 Tuning the Hyperparameters

The efficiency of HMC depends on three hyperparameters: the step size ε , the number of leapfrog steps L , and the mass matrix M . Poor choices lead to either high rejection rates (if ε is too large), slow exploration (if L is too small), or wasted computation (if L is too large and the trajectory doubles back on itself).

The most widely used adaptive scheme is the No-U-Turn Sampler (NUTS) [15], which eliminates the need to specify L by dynamically extending the trajectory until a U-turn criterion is met. While effective in general, NUTS can require many leapfrog steps per sample – each of which, in our setting, demands a full forward and adjoint solve of the electrochemical model.

In this work we instead use the *ChEES-HMC* criterion of Hoffman et al. [14], which adapts ε , L , and M jointly during a warm-up phase by maximising the *Change in the Estimator of the Expected Square* (ChEES) of the trajectory length. We found empirically that ChEES-HMC selected shorter trajectory lengths than NUTS for this problem whilst maintaining comparable effective sample sizes. Since each leapfrog step requires a gradient evaluation – and each gradient evaluation requires a full forward and adjoint solve – reducing the number of integration steps per sample

directly reduces the computational cost. This makes ChEES-HMC particularly well suited to scientific applications in which the gradient of the forward model is expensive to evaluate.

During warm-up, the step size and trajectory length are adapted alongside the mass matrix, and the sampler simultaneously performs gradient descent toward the posterior mode, serving the dual purpose of burn-in and hyperparameter tuning. Once the warm-up phase concludes, the hyperparameters are fixed and all subsequent samples are retained as draws from the posterior.

Chapter 7

Heterogeneous Reaction

We will now be extending the methodology we have built up over the previous chapters to a more complicated electrochemical reaction involving two redox reactions and a heterogeneous chemical reaction that couples the product and the reactant from the two equations which occurs at the surface of the working electrode.



This type of reaction appears in electroreduction of halonitroaromatic compounds in aprotic media [4].

7.1 Mathematical Description

We are able to express the reaction using the same non-dimensional scheme as for the electron reaction. Only that we now need to model the A, B, C and D species in separate governing equations. The electrode boundary conditions change to:

$$\begin{aligned} D_{\text{A}} \left(\frac{\partial c_{\text{A}}}{\partial x} \right)_{x=0} &= k_{\text{red}}^{(1)} c_{\text{A}}(0, t) - k_{\text{ox}}^{(1)} c_{\text{B}}(0, t) \\ D_{\text{B}} \left(\frac{\partial c_{\text{B}}}{\partial x} \right)_{x=0} &= -k_{\text{red}}^{(1)} c_{\text{A}}(0, t) + k_{\text{ox}}^{(1)} c_{\text{B}}(0, t) + k_{\text{het}} c_{\text{B}}(0, t) \\ D_{\text{C}} \left(\frac{\partial c_{\text{C}}}{\partial x} \right)_{x=0} &= k_{\text{red}}^{(2)} c_{\text{C}}(0, t) - k_{\text{ox}}^{(2)} c_{\text{D}}(0, t) - k_{\text{het}} c_{\text{B}}(0, t) \\ D_{\text{D}} \left(\frac{\partial c_{\text{D}}}{\partial x} \right)_{x=0} &= -k_{\text{red}}^{(2)} c_{\text{C}}(0, t) + k_{\text{ox}}^{(2)} c_{\text{D}}(0, t) \end{aligned} \tag{7.2}$$




Figure 7.1:

Using the non-dimensional as we have described before and assuming equal diffusion coefficients:

$$T > 0, X = 0 :$$

$$\begin{aligned} \left(\frac{\partial C_A}{\partial X} \right)_{X=0} &= K_{\text{red}}^{(1)} C_A(0, T) - K_{\text{ox}}^{(1)} C_B(0, T) \\ \left(\frac{\partial C_B}{\partial X} \right)_{X=0} &= -K_{\text{red}}^{(1)} C_A(0, T) + (K_{\text{ox}}^{(1)} + K_{\text{het}}) C_B(0, T) \\ \left(\frac{\partial C_C}{\partial X} \right)_{X=0} &= K_{\text{red}}^{(2)} C_C(0, T) - K_{\text{ox}}^{(2)} C_D(0, T) - K_{\text{het}} C_B(0, T) \\ \left(\frac{\partial C_D}{\partial X} \right)_{X=0} &= -K_{\text{red}}^{(2)} C_C(0, T) + K_{\text{ox}}^{(2)} C_D(0, T) \end{aligned} \quad (7.3)$$

where the dimensionless rate constant of the heterogeneous chemical reaction is given by

$$K_{\text{het}} = \frac{k_{\text{het}} \epsilon}{D_A} \quad (7.4)$$

Notice the K_{het} term in the boundary for both B and C. Regarding the current response the total current will be given by the sum of the contribution of the different heterogeneous *electrochemical* reactions:

$$\frac{I}{FA} = - \left\{ D_A \left(\frac{\partial c_A}{\partial x} \right)_{x=0} + D_C \left(\frac{\partial c_C}{\partial x} \right)_{x=0} + k_{\text{het}} c_B(0, t) \right\} \quad (7.5)$$

Chapter 8

Adsorption Reaction

Chapter 9

Conclusion

Bibliography

- [1] Michael Betancourt. A Conceptual Introduction to Hamiltonian Monte Carlo, July 2018. arXiv:1701.02434 [stat].
- [2] James Bradbury, Roy Frostig, Peter Hawkins, Matthew James Johnson, Chris Leary, Dougal Maclaurin, George Necula, Adam Paszke, Jake VanderPlas, Skye Wanderman-Milne, and Qiao Zhang. JAX: composable transformations of Python+NumPy programs, 2018.
- [3] Haotian Chen, Chenyang Huang, Alexander Rodriguez, Aashutosh Mistry, and Venkatasubramanian Viswanathan. Differentiable electrochemistry: A paradigm characterizing physical laws in electrochemical systems. *ACS Energy Letters*, 11(2):2005–2018, 2026.
- [4] R. G. Compton, Eduardo Laborda, and Kristopher R. Ward. *Understanding voltammetry : simulation of electrode processes*. Imperial College Press, London, 2014 - 2014.
- [5] John Crank. *The Mathematics of Diffusion*. Oxford University Press, Oxford, 2nd edition, 1979.
- [6] Edmund J.F. Dickinson and Andrew J. Wain. The Butler-Volmer equation in electrochemical theory: Origins , value, and practical application. *Journal of Electroanalytical Chemistry*, 872:114145, September 2020.
- [7] A. Einstein. ber die von der molekularkinetischen theorie der wrme geforderte bewegung von in ruhenden flssigkeiten suspendierten teilchen. *Annalen der Physik*, 322(8):549–560, 1905.
- [8] David J. Gavaghan, Jonathan Cooper, Aidan C. Daly, Christopher Gill, Kathryn Gillow, Martin Robinson, Alexandr N. Simonov, Jie Zhang, and Alan M. Bond. Use of bayesian inference for parameter recovery in dc and ac voltammetry. *ChemElectroChem*, 5(6):917–935, 2018.

- [9] Andrew Gelman, John B. Carlin, Hal S. Stern, David B. Dunson, Akti Vehtari, and Donald B. Rubin. *Bayesian Data Analysis*. Chapman & Hall/CRC Texts in Statistical Science Series. CRC, Boca Raton, Florida, third edition, 2013.
- [10] Andrew Gelman, Walter R Gilks, and Gareth O Roberts. Weak convergence and optimal scaling of random walk metropolis algorithms. *The annals of applied probability*, 7(1):110–120, 1997.
- [11] Charles J. Geyer. Introduction to Markov chain Monte Carlo. In Steve Brooks, Andrew Gelman, Galin L. Jones, and Xiao-Li Meng, editors, *Handbook of Markov Chain Monte Carlo*, pages 3–48. CRC Press, Boca Raton, FL, 2011.
- [12] Nikolaus Hansen and Andreas Ostermeier. Completely derandomized self-adaptation in evolution strategies. *Evolutionary computation*, 9(2):159–195, 2001.
- [13] W Keith Hastings. Monte carlo sampling methods using markov chains and their applications. 1970.
- [14] Matthew Hoffman, Alexey Radul, and Pavel Sountsov. An Adaptive-MCMC Scheme for Setting Trajectory Lengths in Hamiltonian Monte Carlo. In *Proceedings of The 24th International Conference on Artificial Intelligence and Statistics*, pages 3907–3915. PMLR, March 2021.
- [15] Matthew D Hoffman, Andrew Gelman, et al. The no-u-turn sampler: adaptively setting path lengths in hamiltonian monte carlo. *J. Mach. Learn. Res.*, 15(1):1593–1623, 2014.
- [16] P Kidger. *On neural differential equations*. PhD Thesis, University of Oxford, 2021.
- [17] Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.
- [18] Graham P Morris, Alexandr N Simonov, Elena A Mashkina, Rafel Bordas, Kathryn Gillow, Ruth E Baker, David J Gavaghan, and Alan M Bond. A comparison of fully automated methods of data analysis and computer assisted heuristic methods in an electrode kinetic study of the pathologically variable [fe (cn) 6] 3–/4–process by ac voltammetry. *Analytical chemistry*, 85(24):11780–11787, 2013.

- [19] Radford M. Neal. *MCMC using Hamiltonian dynamics*. May 2011.
arXiv:1206.1901 [stat].